


Arc Codex — Four-Node Architecture

Role separation settled 2026-09-11 · internal reference, not the public page



syncs when online

Flow — five stages, four nodes, three inter-node hops.

1. **Ingest** — node 1 pulls RSS and lands new articles in the datastore.
2. **Analysis** — node 2 picks up jobs from the datastore, runs the three ARC passes lazily on an article's first view.
3. **Broadcast script** — still on node 2: an additional analysis pass writes an original short piece from the Red/Blue/Purple findings. This piece, not the source article, is what gets read aloud.
4. **Synthesis** — node 3 picks up the broadcast script from the datastore and renders it to mono MP3 with Kokoro.
5. **Publication** — node 1 serves the article, its analysis, and the audio to readers.

Node 4 (portable authoring) is off-pipeline by design: it exists so authoring work can proceed without the fleet reachable, and it syncs back when it rejoins. It never handles live ingest, analysis, synthesis, or publication.

All hops between nodes are mediated by the datastore on node 1; there is no direct node-to-node channel. That is intentional — node 1's availability is already required for the site to serve, so making it the queue substrate adds no new failure surface.